

Set similarity with b-bit k-permutation Minwise Hashing

Simon Wehrli

30.5.2013

1 Introduction

This report is written for the seminar titled “Algorithms for Database Systems” at ETH Zürich. The seminar participants read and summarize various papers, which treat solving problems in the context of Big Data, which is this year’s topic. This report summarizes and explains the core concepts of [1] and [2].

With Big Data, most problems emphasis shifts from single computational complexity to memory space usage considerations. Typically, the runtime is dominated by the time needed for memory accesses, hence the goal is to reduce the number of memory accesses. Often this is achieved by more compact data structures and at the price of less accurate result.

Many today’s applications are faced with very large datasets. A common task is to find *similarity* between two or several such sets. There are lots of problem solutions which use a mapping to sets of properties and then do a *similarity search* on them. Fast (approximative) algorithms for this search enable improvements of many well-known algorithms, e.g. of machine learning and computer vision.

We start by explaining the original MINWISE HASHING, which nicely presents the basic idea behind all algorithms. We move on to a major improvement, the B-BIT K-PERMUTATION MINWISE HASHING, which reduces storage at the cost of more iterations in the algorithm and is a generalization of the concept. We then present ONE PERMUTATION HASHING, which achieves surprising accuracy using only one permutation. We focus on the concepts and algorithms and will reference to papers for applications in practice.

1.1 Notations

We denote by Ω the set of all possible items of the sets $S_n \subseteq \Omega$, $n = 1$ to N . $|\Omega| = D$ is always large (e.g. $D = 2^{64}$). Often we consider only two sets $S_1 = X$, $S_2 = Y$. Let $a = |X \cap Y|$, $b = |X| - a$ and $c = |Y| - a$. Figure 1 on the following page summarizes these definitions graphically.




Figure 1: Two example sets in Ω and the notation a, b, c for the sizes of important subsets.

Similarity There exists different measurements of similarity. We will use the following throughout the paper:

Definition 1. *The normalized similarity between two sets X and Y , known as resemblance or Jaccard similarity, denoted by R , is*

$$R = \frac{|X \cap Y|}{|X \cup Y|} = \frac{a}{|X| + |Y| - a} = \frac{a}{a + b + c} \quad (1)$$

Probability Later we often look at some important event of a problem where we use $\Pr[\cdot]$ as a shorthand for the *Probability* of that event. To get an estimator for R out of a probabilistic argument, we always use some Bernoulli experiment. This is a process, where we repeatedly flip a possibly biased coin, but the bias does not change. We can see it as a sequence of binary random variables, because only two outcomes are possible, 0 or 1. An event is described by an equation. We use the notation $1\{equation\}$ which is one if and only if the *equation* in curly braces is true:

$$1\{equation\} = \begin{cases} 1 & \text{if } equation \text{ evaluates to } true \\ 0 & \text{otherwise} \end{cases}$$

Runtime/Storage Analysis When we do the analysis of an algorithm, we use the big- O -notation. For storage requirement analysis we count in bits, which belongs to the *logarithmic cost model*. We do not count in bytes because we later do optimizations on the bit level and these factors indeed matter for practical purposes. Contrary, for runtime analyses we count in the *uniform cost model*, thus each operation needs one unit of time. We make the realistic assumption that $D \leq 2^{64}$, thus the binary representation always fits a 64-bit integer and operations on any integer up to 64-bit need some constant amount of time.

1.2 Motivating Example

Consider a web search provider, which want to present a result list of web pages without duplicates. To achieve that, for every pair of web pages, we drop one of them, if they are textually very similar. This is the case when their *resemblance* R is greater than some threshold R_0 . But to be able to use the *resemblance* as a measurement, we have to map each page to a set. One could imagine to define this mapping from the page to the set of all words occurring on that page. But with this mapping would not keep track of the order in which their appear. As in several studies [3, 4] we will instead map a page to a set of *shingles*. A shingle is a string of w contiguous words, and we include a *shingle* in our result set if the *shingle* occurs on the page (in the same order). Typically we choose $w = 5$. Figure 2 shows an exemplary mapping.




Figure 2: The mapping from a web page to a set of *shingles*.

Clearly, the number of possible shingles and therefore D is huge. Assuming 10^5 different English words, we have $D = (10^5)^5 \gg 2^{64}$. Thus computing the exact similarities for all pairs of pages of a web search would require prohibitive storage. In general we would already need $\Theta(D)$ storage for a single pair. We need approximation algorithms to improve significantly on this bound.

However, finding duplicates out of m objects, e.g. web pages, needs $O(m^2)$ comparisons even when approximating the similarity. This is a different problem and there are a number of approaches to deal with it, but we will not discuss it here and refer to [4, 5].

2 Original Minwise Hashing

A working approximative solution to this problem was described by Broder and his colleagues [3, 4]. We demonstrate their algorithm on a small running example. Suppose we have two web pages P_X and P_Y and want to compare them:



We map both pages to sets of length-2-shingles. So P_X becomes {“uv”, “vu”} and P_Y becomes {“wv”, “vu”}. Note that the mapping is indifferent to that “uv” occurs twice on P_X . When we continuously number all possible shingles

$\{\text{"uu"}, \text{"uv"}, \text{"uw"}, \text{"vu"}, \text{"vv"}, \text{"vw"}, \text{"wu"}, \text{"wv"}, \text{"ww"}\}$ and define this to be Ω_{shingle} , we can visualize sets in a binary matrix

$$M \in \{0,1\}^{N \times D}, \quad M_{ni} = 1 \{i \in S_n\}. \quad (2)$$

Applied to our example, we get

$$M = \begin{array}{c} \text{Sets:} \\ X: \\ Y: \end{array} \overbrace{\begin{pmatrix} & \text{uu} & \text{uv} & \text{uw} & \text{vu} & \text{vv} & \text{vw} & \text{wu} & \text{wv} & \text{ww} \\ \begin{pmatrix} 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 1 & 0 \end{pmatrix} \end{pmatrix}}^{\Omega_{\text{shingle}}}. \quad (3)$$

Because working with an Ω_{shingle} of strings is nasty, we assume there is a perfect hash function applied to the elements of the original domain which always gives us $\Omega = \{0, 1, \dots, D-1\}$. Additionally we colour the columns of the matrix according to the three coloured subsets in Figure 1 on page 2, e.g. a column with both entries equals to one represents an element contained in both sets. M is rewritten with the new Ω as

$$M = \begin{array}{c} \text{Sets:} \\ X: \\ Y: \end{array} \overbrace{\begin{pmatrix} & 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 \\ \begin{pmatrix} 0 & \textcolor{blue}{1} & 0 & \textcolor{red}{1} & 0 & 0 & 0 & \textcolor{orange}{0} & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & \textcolor{orange}{1} & 0 \end{pmatrix} \end{pmatrix}}^{\Omega}. \quad (4)$$

$\textcolor{blue}{b} \quad \textcolor{red}{a} \quad \textcolor{orange}{c}$

For reasons which later become clear, suppose a random permutation π is performed on Ω ,

$$\pi : \Omega \longrightarrow \Omega.$$

To simplify notation, we overload the definition of π to work also for subsets of Ω . Thus with $\pi(S_n)$ we denote the application of the permutation π to every element of the set S_n . More precisely,

$$\begin{aligned} \pi : 2^\Omega &\longrightarrow 2^\Omega, \\ \pi : S_n &\longmapsto \pi(S_n) = \{\pi(i) | i \in S_n\}. \end{aligned}$$

In the matrix representation, the application of π can be seen as rearranging the matrix columns in a random order. For an exemplary chosen π this could look like

$$M' = \begin{array}{c} \text{Sets:} \\ \pi(X): \\ \pi(Y): \end{array} \overbrace{\begin{pmatrix} & 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 \\ \begin{pmatrix} 0 & 0 & \textcolor{orange}{0} & 0 & 0 & \textcolor{red}{1} & 0 & 0 & \textcolor{blue}{1} \\ 0 & 0 & \textcolor{orange}{1} & 0 & 0 & \textcolor{red}{1} & 0 & 0 & \textcolor{blue}{0} \end{pmatrix} \end{pmatrix}}^{\pi(\Omega)}. \quad (5)$$

$\Delta \quad \Delta \quad \blacktriangle$
 $\textcolor{orange}{c} \quad \textcolor{red}{a} \quad \textcolor{blue}{b}$

Suppose now there is a pointer \blacktriangle starting at the left most column of M' , moving to the right until it sees the first column where there is at least one 1 (i.e. skip columns of only zeros). In our example \blacktriangle would stop at index 2. We can define the event ω , that \blacktriangle stops at a column with two ones, which corresponds to the subset a . We ask for the probability of ω and find

$$\Pr \left[\blacktriangle \text{ stops at } \begin{pmatrix} 1 \\ 1 \end{pmatrix} \right] = \Pr[\omega] = \frac{a}{a+b+c} = R,$$

because we randomly permuted the columns. To determine if or if not ω arises for the given permutation π , we only need the index of the left most 1 in each row respectively the smallest element of each set S_n permuted with π . We define this smallest element of the permuted set¹ as

$$h_{S_n, \pi} = \min(\pi(S_n)).^2 \quad (6)$$

This key observation leads to the crucial

Lemma 1. *For any two sets $X, Y \subseteq \Omega$*

$$\Pr[\min(\pi(X)) = \min(\pi(Y))] = \Pr[h_{X, \pi} = h_{Y, \pi}] = \frac{|X \cap Y|}{|X \cup Y|} = R. \quad (7)$$

In our running example we get

$$\begin{aligned} h_{X, \pi} &= \min(\pi(\{1, 3\})) = \min(\{5, 8\}) = 5, \\ h_{Y, \pi} &= \min(\pi(\{3, 7\})) = \min(\{2, 5\}) = 2. \end{aligned}$$

As we already noticed in the matrix M' , the minima do not agree for our exemplary permutation π .

2.1 The Estimator

To estimate R via a *Bernoulli Process* we must evaluate ω not only once but on k independent random permutations $\pi_1, \pi_2, \dots, \pi_k$. By averaging the outcomes we finally build the estimator

$$\hat{R}_M = \frac{1}{k} \sum_{j=1}^k 1 \{h_{X, \pi_j} = h_{Y, \pi_j}\}, \quad (8)$$

$$\text{Var}(\hat{R}_M) = \frac{1}{k} R(1 - R). \quad (9)$$

It can be proven, that $E(\text{Var}(\hat{R}_M)) = R$ which means that the estimator is unbiased. We later refer to h_{S_n, π_j} as a **sample** and to k as the **sample size**.

¹Note that the smallest element changes under the permutations because π is a permutation on Ω and S_n is only a subset of Ω in general.

²The denotation h is chosen, because the mapping can be interpreted as a hash function with desired properties for *resemblance* estimation.

Many applications, especially duplicate detection tools, are interested in detecting somewhat high similarity, thus an approximation seems reasonable. From (9) on the previous page we learn, that the accuracy can be adjusted by choosing the **sample size** appropriately. In practice we typically have a k between 2^5 to 2^8 to get an absolute error less than 0.01. Even though k can be considered as a constant, the storage requirement of the pre-processing often gets impractical.

2.2 The Algorithm

We have to store a **sample** for each set and this for each permutation. But we can reuse the same k permutations for all sets and pre-compute the k **samples** for each set individually. This can be done in a pre-processing phase. Often we do not care if this phase takes some time, but the space requirements are indeed relevant:

- We need $O(kD \log(D))$ bits of space to store the k permutations in k D -dimensional vectors with entries in the domain Ω .
- Each **sample** potentially needs $O(\log(D))$ bits, because the smallest element can be any element in D . By using k permutations, we need $O(k \log(D))$ bits in total.

In the second phase, we receive queries to estimate the *resemblance* of two specific sets. Here we only have to make the computations from (8) on the preceding page, thus $O(k)$ steps.

Altogether we learn that the algorithm has to consist of two phases, the pre-processing and the *resemblance* estimation. Algorithm 1 presents the procedure of (k -permutation) MINWISE HASHING.

Algorithm 1 Original MINWISE HASHING algorithm, applied to estimating pairwise *resemblances* in a collection of N sets.

Input: Sets $S_n \subseteq \Omega = \{0, 1, \dots, D-1\}, n = 1$ to N . $\triangleright D = |\Omega|$

Output: Estimated *resemblance* \hat{R}_M

// Pre-processing

Generate k random permutations $\pi_j : \Omega \rightarrow \Omega, j = 1$ to k

for all $n = 1$ to $N, j = 1$ to k **do**

Store $\min(\pi_j(S_n))$, denoted by h_{S_n, π_j} .

end for

// Estimation (Use two sets X, Y as an example)

Estimate the *resemblance* by $\hat{R}_M = \frac{1}{k} \sum_{j=1}^k 1 \{h_{X, \pi_j} = h_{Y, \pi_j}\}$

2.3 Deriving the Hamming Distance

Another well-known measurement for the similarity is the *hamming distance*. For the purpose of calculating the *hamming distance* between two sets $X, Y \subseteq$

$\Omega = \{0, 1, \dots, D-1\}$, the sets are first mapped to a D -dimensional binary vector x, y resp.:

Definition 2. Let vector $x, y \in \{0, 1\}^D$, $x_i = 1 \{i \in X\}$, $y_i = 1 \{i \in Y\}$. The hamming distance between X and Y is

$$H = \sum_{i=0}^{D-1} [x_i \neq y_i] = |X \cup Y| - |X \cap Y| = |X| + |Y| - 2a \quad (10)$$

If we reformulate 1 on page 2 as

$$a = \frac{R}{1+R}(|X| + |Y|), \quad (11)$$

we can use the MINWISE HASHING algorithm and any of its improvements in the next sections to estimate H with

$$\hat{H}_b = |X| + |Y| - 2 \frac{\hat{R}_b}{1 + \hat{R}_b}(|X| + |Y|) = \frac{1 - \hat{R}_b}{1 + \hat{R}_b}(|X| + |Y|) \quad (12)$$

Experiments show that the approach via MINWISE HASHING variants is significantly faster than standard methods for computing the *hamming distance*.

3 b-bit k-permutation Minwise Hashing

We will present an algorithm which improves on the storage requirements of the original MINWISE HASHING and is at the same time a generalization of the concept. The idea is to reduce the size of each **sample** by only taking b bits of it, as opposed to, e.g. 64 bits. Intuitively, this will increase the estimation variance $Var(\hat{R}_M)$, at the same **sample size** k . To maintain the same accuracy, we have to increase k . One can show, if the *resemblance* is not too small, we will not have to increase k much and in total use less storage.

For example, when $b = 1$ and $R = 0.5$, the estimation variance will increase at most by a factor of 3. To keep the same accuracy, we have to increase the **sample size** by a factor of 3. If we before stored each minimum $\min(\pi_j(S_n))$ using 64 bits, the improvement with $b = 1$ is $64/3 = 21.3$.

Consider again two sets $X, Y \subseteq \Omega$, on which a random permutation $\pi : \Omega \rightarrow \Omega$ is applied. We extend our notation with

$$h_{X,b,\pi} = b \text{ lowest bits of } \min(\pi(X))$$

$$h_{Y,b,\pi} = b \text{ lowest bits of } \min(\pi(Y))$$

Theorem 1. Assume D is large.

$$P_b = \Pr[1 \{h_{X,b,\pi} = h_{Y,b,\pi}\}] = C_{1,b} + (1 - C_{2,b})R \quad (13)$$

$$r_X = \frac{|X|}{D}, \quad r_Y = \frac{|Y|}{D} \quad (14)$$

$$C_{1,b} = A_{X,b} \frac{r_Y}{r_X + r_Y} + A_{Y,b} \frac{r_X}{r_X + r_Y}, \quad (15)$$

$$C_{2,b} = A_{X,b} \frac{r_X}{r_X + r_Y} + A_{Y,b} \frac{r_Y}{r_X + r_Y}$$

$$A_{X,b} = \frac{r_X [1 - r_X]^{2^b - 1}}{1 - [1 - r_X]^{2^b}}, \quad A_{Y,b} = \frac{r_Y [1 - r_Y]^{2^b - 1}}{1 - [1 - r_Y]^{2^b}} \quad (16)$$

The intuition for the additional terms $C_{1,b}$ and $C_{2,b}$ in (13) compared to (7) on page 5 is that we have to account for a type of “false positive”: When two minima agree on their last b bits, $h_{X,b,\pi} = h_{Y,b,\pi}$, it’s still possible that their are different, $h_{X,\pi} \neq h_{Y,\pi}$. Thus even when $R = 0$, the collision probability P_b is not zero, but rather $C_{1,b}$. This makes the derivation much more complicated.³

Even though Theorem 1 assumes D to be large, experiments show that even for $D = 20$ the absolute error caused by using (13) is less than 0.01.

3.1 The Estimator

From (13) of Theorem 1 we derive the estimator \hat{R}_b for R :

$$\hat{R}_b = \frac{\hat{P}_b - C_{1,b}}{1 - C_{2,b}} \quad (17)$$

$$\hat{P}_b = \frac{1}{k} \sum_{j=1}^k 1 \{h_{X,b,\pi_j} = h_{Y,b,\pi_j}\} \quad (18)$$

This estimator is unbiased, i.e. $E[\hat{R}_b] = R$. Furthermore, the variance of \hat{R}_M converges to the variance of \hat{R}_b , i.e.

$$\lim_{b \rightarrow \inf} \text{Var}(\hat{R}_b) = \frac{R(1 - R)}{k} = \text{Var}(\hat{R}_M) \quad (19)$$

3.2 The Algorithm

Based on the theoretical results, Algorithm 2 presents the procedure of b -bit (k -permutation) MINWISE HASHING.

³A proof of Theorem 1 can be found in the appendix of [6].

Algorithm 2 B-BIT MINWISE HASHING algorithm, applied to estimating pairwise *resemblances* in a collection of N sets.

Input: Sets $S_n \subseteq \Omega = \{0, 1, \dots, D-1\}, n = 1$ to N . $\triangleright D = |\Omega|$

Output: Estimated *resemblance* \hat{R}_b

// Pre-processing

Generate k random permutations $\pi_j : \Omega \rightarrow \Omega, j = 1$ to k

for all $n = 1$ to $N, j = 1$ to k **do**

Store the lowest b bits of $\min(\pi_j(S_n))$, denoted by h_{S_n, b, π_j} .

end for

// Estimation (Use two sets X, Y as an example)

Compute $\hat{P}_b = \frac{1}{k} \sum_{j=1}^k 1 \{h_{X, b, \pi_j} = h_{Y, b, \pi_j}\}$

Estimate the *resemblance* by $\hat{R}_b = \frac{\hat{P}_b - C_{1,b}}{1 - C_{2,b}}$, where $C_{1,b}$ and $C_{2,b}$ are from Theorem 1 on the previous page

3.3 Drawbacks

The major problem of B-BIT MINWISE HASHING is the costly pre-processing. Consider an application in machine learning, where the sets S_n represents some properties of an object (e.g. a document or an image). Often one wants to add objects dynamically at runtime and compare them to other objects by finding the *resemblance* of their properties. Finding the k minima under the permutations may take too long, i.e. uses $O(k|S_n|)$ to $O(kD)$ time per set. In general, the costly pre-processing may cause problems in user-facing applications, where the testing efficiency for new data objects is crucial. There is the need for a entirely fast algorithm to keep this applications responsive.

Another drawback is that storing k permutations is sometimes impractical. If e.g. $D = 10^9$, one permutation vector uses 4GB, which is still possible to store. But the space needed to store e.g. $k = 500$ permutation vectors (each of length D), is not tolerable.

4 One Permutation Hashing

This method is directly motivated by the optimization potential of the standard MINWISE HASHING method: intuitively, it ought to be “wasteful” in that all elements in a set are permuted, scanned but only the minimum will be used. As the name already suggests, we reduce the pre-processing phase to only one permutation.

We will setup a running example with $X, Y, Z \subseteq \Omega = \{0, 1, \dots, 15\}$. Let be π some random permutation on Ω and the already permuted sets be

$$\pi(X) = \{2, 4, 7, 13\}, \quad \pi(Y) = \{0, 3, 6, 13\}, \quad \pi(Z) = \{0, 1, 10, 12\}.$$

As in section 2.3 on page 6 we will represent sets $S_n \subseteq \Omega$ as vectors $s_n \in \{0, 1\}^D$, $(s_n)_i = 1 \{i \in S_n\}$. Sticking them together as rows of a matrix, we get

a binary matrix M equally to the one defined in (2) on page 4. We show the matrix M' where the columns are already permuted by π :

$$M' = \begin{array}{l} \text{Sets:} \\ \pi(X): \\ \pi(Y): \\ \pi(Z): \end{array} \overbrace{\begin{pmatrix} 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 & 10 & 11 & 12 & 13 & 14 & 15 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 \end{pmatrix}}^{\Omega} \quad (20)$$

The idea is to divide the columns evenly into t (here $t = 4$) bins (parts) and take the minimum in each bin. Because later we only compare minima within one bin, we can re-index the elements to use the smallest possible representation:

$$M' = \begin{array}{l} \text{Sets:} \\ \pi(X): \\ \pi(Y): \\ \pi(Z): \end{array} \begin{array}{c} \overbrace{\begin{pmatrix} 0 & 1 & 2 & 3 & 0 & 1 & 2 & 3 & 0 & 1 & 2 & 3 & 0 & 1 & 2 & 3 \\ 0 & 0 & \textcircled{1} & 0 & \textcircled{1} & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & \textcircled{1} & 0 & 0 \\ \textcircled{1} & 0 & 0 & 1 & 0 & 0 & \textcircled{1} & 0 & 0 & 0 & 0 & 0 & 0 & \textcircled{1} & 0 & 0 \\ \textcircled{1} & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & \textcircled{1} & 0 & \textcircled{1} & 0 & 0 & 0 & 0 \end{pmatrix}}^{\Omega'} \\ \text{bin 1} & \text{bin 2} & \text{bin 3} & \text{bin 4} \end{array} \quad (21)$$

We get the minima-vectors

$$\begin{aligned} v_X &= [2, 0, *, 1], \\ v_Y &= [0, 2, *, 1], \\ v_Z &= [0, *, 2, 0], \end{aligned} \quad (22)$$

where '*' denotes an empty bin.

To derive the *resemblance* between two sets, e.g. X, Y , we introduce two definitions:

$$\begin{aligned} \text{number of "jointly empty bins":} \quad N_{emp} &= \sum_{j=1}^t I_{emp,j}, \\ \text{number of "matched bins":} \quad N_{mat} &= \sum_{j=1}^t I_{mat,j}, \end{aligned} \quad (23)$$

where $I_{emp,j}$ and $I_{mat,j}$ are defined for the j -th bin, as

$$\begin{aligned} I_{emp,j} &= \begin{cases} 1 & \text{if both } \pi(X) \text{ and } \pi(Y) \text{ are empty in the } j\text{-th bin} \\ 0 & \text{otherwise} \end{cases} \\ I_{mat,j} &= \begin{cases} 1 & \text{if both } \pi(X) \text{ and } \pi(Y) \text{ are not empty and the smallest} \\ & \text{elements in the } j\text{-th bin matches, i.e. } (v_X)_j = (v_Y)_j \\ 0 & \text{otherwise} \end{cases} \end{aligned} \quad (24)$$

4.1 The Estimator

Recall the notations $a = |X \cap Y|$ and $f = |X \cup Y| = |X| + |Y| - a$. We formulate the estimator as

Lemma 2. *The resemblance is estimated by*

$$\hat{R}_{mat} = \frac{N_{mat}}{k - N_{emp}} \quad (25)$$

is unbiased, i.e.

$$\mathbb{E} [\hat{R}_{mat}] = R. \quad (26)$$

In our example we have $N_{emp} = 1$ and $N_{mat} = 1$. Thus $\hat{R}_{mat} = 1/3$.

Because it is a bit surprising that the estimator is unbiased, we give a

Proof of Lemma 2. Because we assume that the sets and thus the data vectors are not completely empty, it holds

$$t - N_{emp} > 0 \Rightarrow P[t - N_{emp} > 0] = 1,$$

hence we get rid of division-by-zero problems and $m > 0$ in (29) is always true. From the definitions in (24) on page 10 follows

$$I_{emp,j} = 1 \Rightarrow I_{mat,j} = 0, \quad (27)$$

$$\mathbb{E} [I_{mat,j} | I_{emp,j} = 0] = R, \quad (28)$$

$$\text{for } m > 0 : \quad \mathbb{E} [I_{mat,j} | t - N_{emp} = m] = \frac{m}{t} R, \quad (29)$$

and we derive

$$\begin{aligned} \mathbb{E} [N_{mat} | t - N_{emp} = m] &\stackrel{(27)}{=} \sum_{j=1}^{t-N_{emp}} \mathbb{E} [I_{mat,j} | t - N_{emp} = m] \\ &\stackrel{(28)}{=} \sum_{j=1}^{t-N_{emp}} R = R(t - N_{emp}). \end{aligned} \quad (30)$$

Finally,

$$\begin{aligned} \mathbb{E} \left[\frac{N_{mat}}{t - N_{emp}} \middle| t - N_{emp} = m \right] &\stackrel{(30)}{=} R, \quad (\text{independent of } N_{emp}) \\ \Rightarrow \mathbb{E} \left[\frac{N_{mat}}{t - N_{emp}} \right] &= \mathbb{E} [\hat{R}_{mat}] = R \end{aligned} \quad (31)$$

□

We give the variance without a proof.⁴

$$\text{Var} [\hat{R}_{mat}] = R(1 - R) \left(\mathbb{E} \left[\frac{1}{t - N_{emp}} \right] \left(1 + \frac{1}{f - 1} \right) - \frac{1}{f - 1} \right) \quad (32)$$

⁴A proof can be found in the appendix of [2].

If we have very few empty bins, i.e. N_{emp} is essentially zero and $t \ll f$, the term simplifies to

$$\text{Var} [\hat{R}_{mat}] \approx \frac{R(1-R)}{t} \left(\frac{f-t}{f-1} \right) \approx \frac{R(1-R)}{t} \quad (33)$$

as expected.

4.2 The Algorithm

Based on the theoretical results, Algorithm 3 presents the procedure of ONE PERMUTATION HASHING.

Algorithm 3 ONE PERMUTATION HASHING algorithm, applied to estimating pairwise *resemblances* in a collection of N sets.

Input: Sets $S_n \subseteq \Omega = \{0, 1, \dots, D-1\}, n = 1$ to N . $\triangleright D = |\Omega|$

Output: Estimated *resemblance* \hat{R}_{mat}

// Pre-processing

Generate one random permutations $\pi : \Omega \rightarrow \Omega$

for all $n = 1$ to N **do**

Permute set S_n with π , store the re-indexed minima of each of the t bins in a data vector $v_{S_n} \in \{0, \dots, \lfloor |S_n|/t \rfloor\}^t$. \triangleright see (21)

end for

// Estimation (Use two sets X, Y as an example)

Estimate the *resemblance* by $\hat{R}_{mat} = \frac{N_{mat}}{t - N_{emp}}$ \triangleright see (23)

Empirical results show that the ONE PERMUTATION HASHING scheme performs as well or even slightly better than the B-BIT K-PERMUTATION HASHING scheme.

Let's have a look at the runtime and storage requirements and compare them with Algorithm 2 on page 9. The pre-processing phase generates one permutation vector, which uses $O(D \log(D))$ storage, which even for large D is still practical (in contrast to storing k permutation vectors for the MINWISE HASHING method). The pre-processing also generates a data vector with t elements (which uses themselves only $O(\log(D/t))$ bits because we re-indexed them within the bins) for each set, thus $O(t \log(D/t))$ bits per set.

Table 1 on page 13 summarizes the runtime and storage requirements of the different algorithms.

Space used for . . .	storing permutation(s)	storing pre-processing data per set
Naïve approach (store whole sets)	—	$O(D)$
original MINWISE HASHING	$O(kD \log(D))$	$O(k \log(D))$
B-BIT MINWISE HASHING	$O(kD \log(D))$	$O(kb)$
ONE PERMUTATION HASHING	$O(D \log(D))$	$O(t \log(D/t))$

Table 1: Algorithm comparison. We only list the space complexity because the processing time boundaries are given by the time needed to read the stored data at least once and hence is dominated by the space complexity.

5 Conclusions

MINWISE HASHING and ONE PERMUTATION HASHING are standard techniques for efficiently estimating set *resemblance* in massive datasets. We started at explaining the original MINWISE HASHING and demonstrated how reducing the amount of stored bits per **sample** to b bits leads to a effective reduction of the required storage and consequently computational overhead.

We then went over to ONE PERMUTATION HASHING, which uses the same basic idea, but makes better use of the permuted set and therefore only needs one permutation. Overall, this last algorithm is superior to the others in many aspects.

References

- [1] Ping Li 0001 and Arnd Christian König. Theory and applications of b -bit minwise hashing. *Commun. ACM*, 54(8):101–109, 2011.
- [2] Ping Li 0001, Art B. Owen, and Cun-Hui Zhang. One permutation hashing for efficient search and learning. *CoRR*, abs/1208.1259, 2012.
- [3] Andrei Z. Broder. On the resemblance and containment of documents. In *Compression and Complexity of Sequences, IEEE Computer Society*, pages 21–29, Italy, 1997. Positano.
- [4] Andrei Z. Broder, Steven C. Glassman, Mark S. Manasse, and Geoffrey Zweig. Syntactic clustering of the web. *Computer Networks*, pages 1157–1166, 1997.
- [5] Moses Charikar. Similarity estimation techniques from rounding algorithms. In *Proceedings of the 34th Annual ACM Symposium on Theory of Computing (STOC-02)*, pages 380–388, New York, May 19–21 2002. ACM Press.
- [6] Ping Li 0001 and Arnd Christian König. b -bit minwise hashing. *CoRR*, abs/0910.3349, 2009.